


The target-relevant action shift stays small until all 512 image positions are replaced

$\pi 0.5 \cdot$ image-position K/V at layers 12–17 $\cdot$ 8 directed prompt-pair cells $\times$ 3 initial states; same pixels, state, and action noise within each comparison.

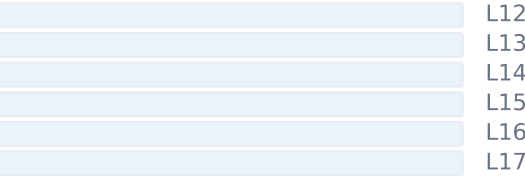
a



b

What each dose changed

K and V at every one of six layers



128/512 object-centered



256/512 object-centered



256/512 random



512/512 all positions



At 512 positions, object-centered and random are the same intervention.

Bold curves: descriptive median over 24 scene-direction rows. Hollow points: eight directed prompt-pair medians (three initial states each). R is target-axis progress, not full-vector similarity; the intervention edits positions, not individual neurons.